

Strategic Defense of Evolutionary Trajectory Fine-Tuning (ETFT)

Y Combinator Interview Argumentation Guide

Lukas Benda

lukas.benda@boldpivot.cz

BoldPivot

April 2026

Abstract

This document outlines the core business trajectory, go-to-market strategy, and defensive arguments for the Evolutionary Trajectory Fine-Tuning (ETFT) framework. It is specifically structured to address high-level investor scrutiny regarding market positioning, data scalability, initial revenue generation, and technical traction.

1 The Frontier Lab Dynamic: Positioning Against Incumbents

The core question: *"If this is so powerful, why wouldn't OpenAI, DeepMind, or Anthropic just do it? Don't they have more compute?"*

1.1 Argumentation: The Symbiotic Licensing Model

We are not competing with frontier labs; we are building a higher-order solution that depends on their base models. The frontier labs are inherently trapped in the "Scaling Law Paradigm" their primary objective is achieving generalized artificial intelligence (AGI) by exponentially increasing parameter counts and compute clusters. They are incentivized to go "up and massive."

ETFT, conversely, is an optimization engine built for extreme constraints. We go "down and dense." Our business trajectory relies on a Symbiotic Licensing Model. We supply frontier labs, research institutions, and sovereign compute centers with the architectural innovation engine they need when sheer parameter scaling hits physical or economic walls. We are the ecosystem's structural engineers, not their competitors.

2 The Data Bottleneck: Scaling the Meta-Dataset ($\mathcal{M}$)

The core question: *"Creating the 'fossil record' requires expert intervention. How do you scale this without an army of PhDs?"*

2.1 Argumentation: The Synthetic Fossil Record

The creation of the evolutionary trajectory is not a manual, human-labor-intensive process. Algorithmic regressions systematically stripping away optimizations while maintaining compilability is a fully deterministic and information-complete task.

The primary workload is carried by AI coding agents (e.g., OpenAI Codex). We utilize these agents to brute-force the regression chains: degrading SOTA implementations down to primordial baselines, verifying operational integrity via our compilation stage-gates, and automatically

logging the resource deltas (Δ Loss, Δ Memory). Human expert intervention is strictly reserved for verifying the final SOTA+1 leaps, allowing us to build a proprietary, highly scalable data moat at a fraction of traditional R&D costs.

3 Go-to-Market: The "First Dollar" Strategy

The core question: *"Who is the very first person to write you a check for this, and what exactly are they buying?"*

3.1 Argumentation: Extreme Constraint Niches

Our initial revenue does not rely on selling generalized LLMs. We target organizations facing the hardest, most lucrative problem domains regarding algorithmic constraints:

- **Edge Computing & Mobile Hardware (e.g., Apple, Qualcomm):** These corporations are desperate to fit functional, highly capable LLMs entirely onto local silicon without draining battery life or exceeding strict thermal limits.
- **Quantitative Finance (HFT):** Funds where algorithmic architectures must be optimized to shave microseconds off execution times, constrained by severe hardware limits.

The Product: These clients are not simply buying the ETFT software; they are licensing the generated **SOTA+1 artifacts** optimized specifically for their hardware, or paying for API access to our Calibration Engine to run against their proprietary codebases.

4 Traction & Proof of Concept: The Stage-Gate Validation

The core question: *"What have you actually built, and how do you prove it works before hypothesizing the future?"*

4.1 Argumentation: The OpenAI Parameter Golf Challenge

We are currently building the system as a reference implementation targeted directly at the OpenAI Parameter Golf challenge. This provides an objective, highly visible, third-party validation crucible.

By applying our Evolutionary Replay mechanism to this specific challenge, we establish a hard algorithmic stage-gate. We discard any SOTA+1 attempts until the system confidently reconstructs the entire trajectory from a dense baseline up to the 16 MB, 10-minute training limit. The Parameter Golf challenge acts as our "Apollo Mission" proving to investors and future enterprise clients that ETFT can systematically deduce architectural breakthroughs under the industry's most brutal constraints.